Attention and Transformers

ECON 8310: Business Forecasting — Lecture 9, Part 2

Department of Economics

University of Nebraska at Omaha

October 29, 2026

Lecture Outline

1 What Recurrence Could Not Do

2 The Attention Mechanism

3 The Transformer Encoder

4 Transformers for Time Series

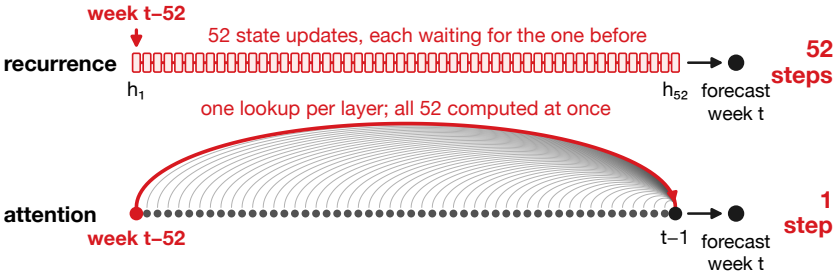
5 Key Takeaways

What Recurrence Could Not Do

Two constraints that are built into reading a sequence one step at a time.

Two Problems Built Into Recurrence

Part 1's forecaster reads strictly left to right. Follow one week, $t - 52$, to the forecast:



The information bottleneck: week $t - 52$ must survive 52 squeezes through one 64-number state.

The sequential bottleneck: each update waits for the last, however many GPUs you own.

Attention removes both at once: every week reads every other directly, and all are computed together. Path length counts operations, not what a trained model uses.

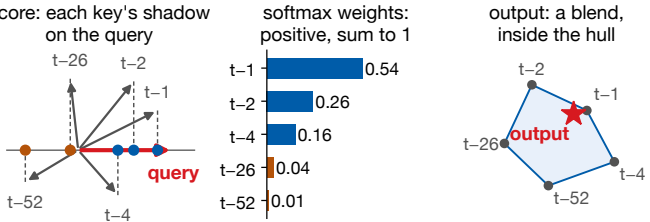
The Attention Mechanism

A learned, weighted lookup — which is a more useful way to read it than the formula suggests.

Scaled Dot-Product Attention (Vaswani et al. 2017)

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right) V$$

Each week's **query** is scored against every week's **key**, and the scores weight the **values**; Q, K, V are learned projections of the same weeks, d_k the key length.

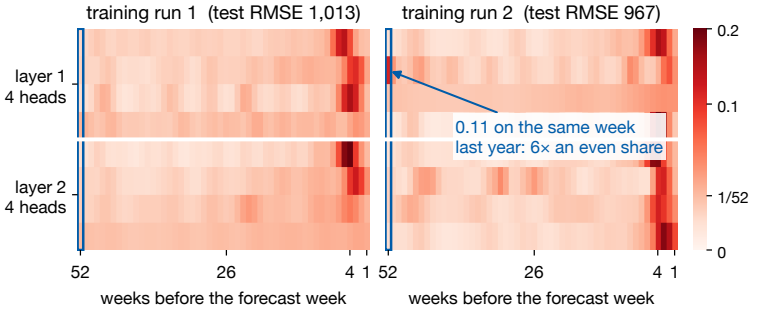


a two-dimensional cartoon of one 16-dimensional head

The score is the query's length times the key's shadow on it. Softmax turns the scores into positive weights that sum to one, so the output stays *inside* the values. Attention averages; it cannot extrapolate. The $\sqrt{d_k}$ keeps the softmax from saturating.

What That Buys a Forecaster

Attention can weight any week in the window in one step; an RNN needs 52 updates to reach last year. Where our trained model actually looked, over the test year:



Mostly the last two to seven weeks; in one of two training runs, one head found last year's week. The weights are **learned, not specified** — and they show where the model looked, not what mattered.

Multi-Head Attention

One attention operation gives one set of weights — one notion of what is relevant.
Multi-head attention runs h of them in parallel on different learned projections, then concatenates:

$$\text{MultiHead}(Q, K, V) = \text{concat}(\text{head}_1, \ldots, \text{head}_h)\, W^O$$

Each head can specialize. Ours did: most watch the last few weeks, one spreads its weight evenly, and in one run one found the annual lag.

The cost is a surprise. Each head works in d_{model}/h dimensions, so eight heads cost about the same as one full-width head.

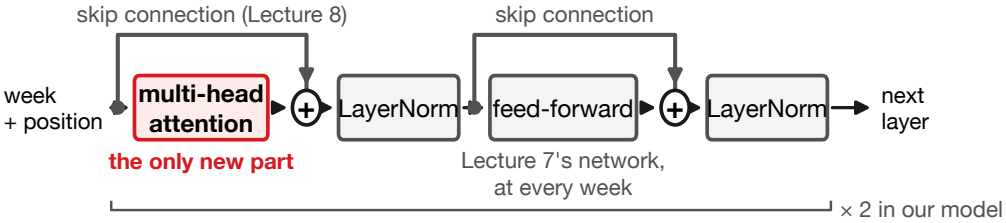
Our model: `nhead=4, d_model=64` — four heads of 16 dimensions.

The Transformer Encoder

Attention, plus the three pieces that make it trainable.

The Encoder Block

A Transformer encoder layer is four components in a fixed arrangement:



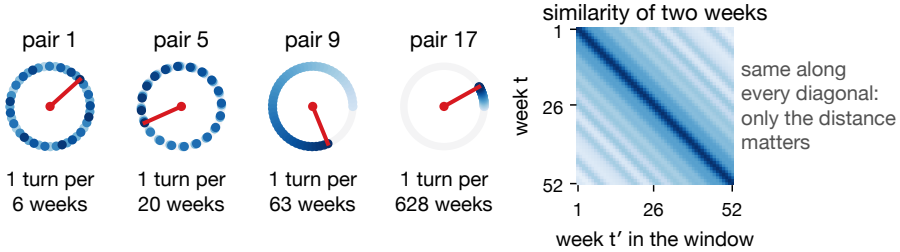
Stack N of those and you have the encoder. Nothing in the block is new to you except attention itself — the skip connections are ResNet's, LayerNorm rescales each week's vector, and the feed-forward sublayer is the Lecture 7 network applied to every week.

One thing is conspicuously missing, and it is not an oversight. Nothing above depends on the *order* of the input.

Positional Encoding: Attention Cannot Tell Time

Attention ignores order: shuffle the window and each week's output moves with it; the rest is an unordered bag. A forecaster that cannot tell two weeks ago from thirty is not forecasting.

The fix adds 32 pairs $(\sin \omega_i t,\ \cos \omega_i t)$, $\omega_i = 10000^{-2i/d}$, to each week — each pair a clock hand turning at its own speed:



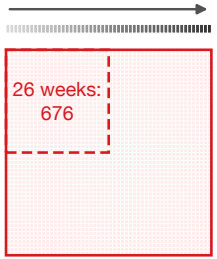
Shifting Δ weeks turns each hand by a fixed angle, so similarity depends only on distance. **We measure what removing it costs at the end.**

Why Transformers Took Over

Parallelism. No score waits for another. More than accuracy, this is why it won.

Direct long-range paths. One step, any distance.

Scale. It kept improving as parameters and data grew. BERT (Devlin et al. 2019) and GPT are this block, repeated.



recurrence: 52 steps, in order

attention: $52 \times 52 = 2{,}704$ scores, all at once

double the window: 2× the steps, 4× the scores

a cell counts an operation, not its cost

Each cell is one query–key score: $O(T^2)$ time and memory, fine at $T = 52$, painful at $T = 10{,}000$.

> Every advantage here is about **large data and long sequences**. Hold that against a 52-week window and 5,160 training windows.

Transformers for Time Series

The implementation, the measurement, and an honest reading of it.

Transformer Encoder in PyTorch

Minimal forecasting encoder

```
layer = nn.TransformerEncoderLayer(d_model=64, nhead=4,
       dim_feedforward=128, dropout=0.1, batch_first=True)
self.enc = nn.TransformerEncoder(layer, num_layers=2)
self.inp = nn.Linear(n_features, 64)
self.fc  = nn.Linear(64, 1)

def forward(self, x):
  h = self.inp(x) + self.pe     # position must be added
  return self.fc(self.enc(h)[:, -1, :]).squeeze(-1)
```

`nn.TransformerEncoderLayer` gives you the whole block — but not the positional encoding. You add that yourself, and forgetting is a silent failure.

Multi-step forecasting also needs a causal mask.

Does It Work?

Same weekly panel and test block as everything since Homework 3; every sequence model reads 52-week windows (Part 1's headline 842 and 987 were at 26 weeks).

Model	Test RMSE	Parameters
LASSO, 46 engineered features	**744**	46
XGBoost	781	—
Vanilla RNN	**805**	4,545
Random forest	899	—
Transformer encoder, 2 layers	**990**	67,329
LSTM	998	17,985
1D CNN, best variant (26-week window)	1,032	8,161
Seasonal naive	2,152	0

The Transformer beats the CNN and **ties the LSTM** — 990 against 998, a difference far smaller than the run-to-run spread. It loses to a vanilla RNN holding **one-fifteenth** as many parameters, and to a 46-coefficient linear model.

Reading That Result Honestly

Every advantage on the "why Transformers took over" slide is an advantage *at scale*, and this problem has none of the scale.

Parallel computation	Real, but irrelevant — the RNN trains in seconds either way.
Direct long-range paths	Worth little when the signal is dominated by recent weeks.
Scales with data	5,160 training windows is not the regime where that pays.

And positional encoding is doing real work, which we can measure. Remove it and the model still knows the last week, where it reads out, but sees the other 51 as an unordered bag:

Transformer **with** positional encoding	**990**
Transformer **without** positional encoding	1,282

292 RMSE — dropping it below even the 1D CNN. That gap is what the order of those 51 weeks is worth, and the reason the sinusoids exist.

Key Takeaways

Six architectures, one series, and what the ranking actually taught us.

Sequence Architecture Comparison

Architecture	Long range	Parallel	Needs
1D CNN	receptive field only	yes	modest data
Vanilla RNN	$\sim$10–30 steps	no	modest data
LSTM	hundreds of steps	no	more data
Transformer	unlimited, $O(T^2)$	**yes**	**much more data**

Choose by the structure of your problem, not by recency of the architecture. Short window and recent-lag signal: a plain RNN, or honestly a linear model. Long sequences, abundant data, dependencies at arbitrary distance: a Transformer, and nothing else comes close.

That second case is language, which is why this architecture reshaped that field and has been slower to displace gradient boosting in tabular forecasting — where the M5 competition was won by gradient-boosted trees, Lecture 5's method.

Lecture 9, Part 2: Key Takeaways

1. **Attention**: score queries against keys, softmax, average the values.
2. Its price is **quadratic**: fine at 52 steps, the central constraint at 10,000.
3. **Multi-head**: several notions of relevance for about the cost of one.
4. An encoder block: attention, **skips**, LayerNorm, feed-forward. Only attention is new.
5. **Attention ignores order**: add position yourself; PyTorch will not.
6. Here it **tied the LSTM** and lost to an RNN **one-fifteenth** its size: it needs scale.

Next: Bayesian Statistics I — a different answer to what a forecast should tell you.

References I

- Devlin, Jacob et al. (2019). "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding". In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*, pp. 4171–4186.
- Vaswani, Ashish et al. (2017). "Attention Is All You Need". In: *Advances in Neural Information Processing Systems*. Vol. 30, pp. 5998–6008.