
Synthesis: What the Results Taught Us

ECON 8310: Business Forecasting — Lecture 13

Department of Economics

University of Nebraska at Omaha

August 25, 2026

Lecture Outline

- 1 What We Built
- 2 What the Results Taught
- 3 Choosing a Method
- 4 Doing It Honestly
- 5 Where to Go Next

What We Built

Twelve weeks, four families of method, and one dataset carried through all of them.

The Arc of the Course

Part	Lectures	The central idea
Time series	1–3	Specify a structure — level, trend, season — and estimate it.
Trees	4–6	Stop specifying. Hand an algorithm features and let it find the structure.
Deep learning	7–9	Let the <i>architecture</i> encode what kind of structure to expect.
Bayesian	10–12	Change the deliverable: return a distribution, not a number.

Read down that last column and the course is not really four toolkits. It is four answers to a single question: **where does the knowledge in a forecast come from?**

From you, in the model structure. From the data, via a search procedure. From the architecture's assumptions. Or from a prior, stated openly and updated. **That framing is what survives when the libraries change.**

The Scoreboard

Every model in this course, on the M5 weekly panel — same 30 series, same held-out block.

Model	Test RMSE	Where
LASSO, 46 engineered features	744	Homework 4
XGBoost, tuned	781	Homework 4
Vanilla RNN, 4 raw channels	842	Lecture 9 Part 1
Random forest	899	Homework 3
LSTM	987	Lecture 9 Part 1
Transformer encoder	990	Lecture 9 Part 2
1D CNN, flatten head	1,032	Lecture 8
Seasonal naive	2,152	every week

The ordering is not the one anyone predicts in Week 1. A 46-coefficient linear model is at the top. The two architectures with the most attention behind them — the LSTM and the Transformer — finished fifth and sixth of eight, behind a plain RNN with a fraction of their parameters.

What the Results Taught

Three lessons that came out of measurements, not out of anyone's opinion.

Lesson 1: The Benchmark Is a Real Competitor

In Homework 1 the seasonal-naive benchmark **beat every exponential smoothing and ARIMA model** you fitted. That was not a trick question — it is what the data supports at that horizon.

Across the whole course the benchmark was beaten, but rarely by the margin the effort implied. The Bayesian structural model beat it by **4%**. The GAMs in Homework 2 beat it by **12%**. The best model on the weekly panel beat it by 65% — after 46 engineered features and a cross-validated penalty.

Compute the benchmark first, every time. It costs one line, it is often close, and it is the only number that tells you whether the rest of the work was worth doing.

This is also the cheapest professional habit you can carry out of this course. A forecast presented without a benchmark is a number no one can evaluate — including you.

Lesson 2: More Sophisticated Did Not Mean Better

Four times this course, the more advanced method lost — and each time for a reason we could name.

LSTM lost to a plain RNN	At three window lengths. Gating fixes long-range forgetting; a 26-week window does not suffer from it. Four times the parameters, no gain.
Transformer lost to an RNN	With $15\times$ the parameters. Its advantages are advantages <i>at scale</i> , and 5,940 windows is not scale.
CNNs lost to LASSO	A seven-week receptive field cannot see an annual lag.
Tuning made XGBoost worse	Nine fits, and the defaults were three RMSE better.

None of these means the losing method is bad — each solves a problem this data does not have. LSTMs earn their gates on long dependency chains; Transformers earn their parameters on large corpora.

Match the method to the structure of the problem, not to the recency of the paper.

Lesson 3: One Idea Appeared Six Times

Lecture	Name it went by	What it penalized
3	GAM smoothing penalty	Wiggleness of a spline
4	Cost-complexity pruning	Number of leaves
5	reg_lambda, gamma	Leaf weights, leaf count
6	Ridge, LASSO, Elastic Net	Coefficient size
7	Weight decay, dropout	Network weights
11	Hierarchical prior	Distance from the group mean

Six model families that look nothing alike, one bargain underneath: **fit against complexity, with a single parameter setting the exchange rate.**

Five of those six tune that parameter by cross-validation. The sixth — the hierarchical prior — **estimates it**, which is the cleanest version of the idea in the course.

This is what to remember when the tools change. The libraries in this course will be replaced. The trade between fit and complexity will not.

Choosing a Method

A decision guide built from what we measured.

Which Method, When

If your situation is...	start here
One series, clear seasonality, short history	ETS or ARIMA — and check the benchmark
Several seasonal cycles at once	GAM or Prophet
Many engineered features, tabular	Regularized regression first , then XGBoost
Many parallel series, shared drivers	Gradient boosting with series identifiers
Genuinely long dependencies, lots of data	Recurrent or attention architectures
Few observations per group	Hierarchical Bayesian — partial pooling
The decision needs a probability	Bayesian, for the predictive distribution

“Regularized regression first” is not a concession — it won our largest comparison and its coefficients can be read aloud. And the last row is the only one where the *question* changes rather than the accuracy.

Doing It Honestly

The errors in this field do not announce themselves.

The Errors That Fail Silently

Every one of these produces a confident, plausible, wrong answer with no error message.

KFold on a time series	Random folds put the future in training. CV error looks excellent; the forecast fails.
Scaling outside the fold	Validation statistics leak into training. Smaller effect, same class of error.
Rolling features unshifted	A window including the current week leaks the target into its own predictor.
Reading MDI importance	Pure noise ranked 17th of 29 in Homework 3 — above a real feature — purely for having many distinct values.
Ignoring $\hat{R}$ and divergences	Our first two Bayesian fits returned reasonable-looking forecasts while failing convergence.

The common thread: the failure mode is not a crash, it is a number that looks fine — which is why the discipline is procedural. Benchmark first, split by time, check the diagnostic, plant a control.

What a Defensible Forecast Report Contains

A benchmark	Computed, not cited. Every model measured against it.
A time-aware split	Walk-forward or a held-out final block, never random.
More than one evaluation	A single test split can flatter. Homework 4 ranked models the same way twice, which is why that ranking is believable.
The result that went against you	Including tuning that did not help, and the model you expected to win.
An interpretation	In the units of the business, not of the loss function.

The fourth row is the one that distinguishes professional work. A report where every result confirms the author's expectation is not a report anyone should act on.

You have practised this all semester: the assignments were built so the obvious answer is frequently the wrong one, and the credit was for noticing.

Where to Go Next

What we left out, and what to read.

What This Course Did Not Cover

Hierarchical reconciliation	Making store forecasts add up to regional totals. Different from Lecture 11's hierarchy. FPP Ch. 11.
Intermittent demand	Series that are mostly zeros — Croston's method and its relatives. Common in spare parts and slow-moving SKUs.
Structural breaks	When the future stops resembling the past. Every method here fails together, and detecting it early is its own field.
Causal inference	DAGs told us confounders exist. Estimating causal effects credibly is a separate course.
Foundation models	Pretrained forecasters applied zero-shot. Moving quickly — FPP-Py Ch. 15 is a starting point.

The third row is the one most likely to matter in your career. Every model in this course assumes the future resembles the past. When that stops being true, the model does not tell you — and the benchmark fails at the same moment.

1. **Compute the benchmark first.** It is often close, and it is the only number that says whether your work was worth doing.
2. **Match the method to the problem,** not to the paper's publication date. The most advanced model on our data finished sixth.
3. **One idea recurs:** fit against complexity, with a parameter setting the rate. You have now seen it six times under six names.
4. **Split by time, always.** The leakage errors in this field do not crash — they return a better-looking number.
5. **Report what went against you.** It is the part that makes the rest credible.
6. **Know what your forecast is for.** A point estimate and a predictive distribution answer different questions, and only one of them supports a decision under uncertainty.

The libraries will change. The habits are the transferable part.

References I
